

By: Yousef Khalil

```
import kagglehub
# Download latest version
path = kagglehub.dataset_download("jp797498e/twitter-entity-sentiment-
analysis")
print("Path to dataset files:", path)

Downloading from
https://www.kaggle.com/api/v1/datasets/download/jp797498e/twitter-
entity-sentiment-analysis?dataset_version_number=2...

100%|██████████| 1.99M/1.99M [00:00<00:00, 99.4MB/s]

Extracting files...
Path to dataset files:
/root/.cache/kagglehub/datasets/jp797498e/twitter-entity-sentiment-
analysis/versions/2

import pandas as pd
import numpy as np
import os
import re
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")

dataset=pd.read_csv(os.path.join(path,os.listdir(path)[1]))
dataset.head()

{"summary": {"name": "dataset", "rows": 74681, "fields": [{"column": "2401", "properties": {"dtype": "number", "std": 3740, "min": 1, "max": 13200, "num_unique_values": 12447, "samples": [1616, 2660, 2335]}, "semantic_type": "", "description": ""}, {"column": "Borderlands", "properties": {"dtype": "category", "num_unique_values": 32, "samples": ["Cyberpunk2077", "Microsoft", "TomClancysRainbowSix"]}, "semantic_type": "", "description": ""}, {"column": "Positive", "properties": {"category": "Positive", "num_unique_values": 4, "samples": ["Neutral", "Irrelevant", "Positive"]}, "semantic_type": "", "description": ""}], "properties": {"column": "2401", "dtype": "number", "min": 1, "max": 13200, "num_unique_values": 12447, "samples": [1616, 2660, 2335], "std": 3740}, {"column": "Borderlands", "dtype": "category", "num_unique_values": 32, "samples": ["Cyberpunk2077", "Microsoft", "TomClancysRainbowSix"]}, {"column": "Positive", "category": "Positive", "num_unique_values": 4, "samples": ["Neutral", "Irrelevant", "Positive"]}}}
```

```

\"description\": \"\\n      }\\n    },\\n    {\\n      \"column\": \"im
getting on borderlands and i will murder you all ,\\\",\\n
\"properties\": {\\n        \"dtype\": \"string\",\\n
\"num_unique_values\": 69490,\\n        \"samples\": [\\n
how does my stained glass open facebook account girl already have 200
likes!!!! and i sure am so!!!! oh thankful!?!?!\\\",\\n          \"How
not to get bored about every damn thing in life.\",\\n          \"The
Best Perfect Way to Protect All the Planet Samsung Galaxy Note10 + By
buff. ly / The 2zkjIhu..\\n        ],\\n        \"semantic_type\":
\"\",\\n        \"description\": \"\\n      }\\n    }\\n  ]\\
n}\",\"type\":\"dataframe\",\"variable_name\":\"dataset\"}

dataset.rename(columns={dataset.columns[3]:"reviews",dataset.columns[0]
]: "index",dataset.columns[2]:"label",dataset.columns[1]:"company"},inplace=True)

dataset.sample(3)

{"summary": {"\n  \"name\": \"dataset\",\\n  \"rows\": 3,\\n  \"fields\":
[\n    {\n      \"column\": \"index\",\\n      \"properties\": {\\n
\"dtype\": \"number\",\\n        \"std\": 3720,\\n        \"min\":
1386,\\n        \"max\": 8809,\\n        \"num_unique_values\": 3,\\n
\"samples\": [\\n          8809,\\n          5544,\\n          1386\\n
],\\n        \"semantic_type\": \"\",\\n        \"description\": \"\\n
\"},\\n      {\\n        \"column\": \"company\",\\n        \"properties\": {\\n
\"dtype\": \"string\",\\n        \"num_unique_values\": 3,\\n
\"samples\": [\\n          \"Nvidia\",\\n          \"Hearthstone\",\\n          \"Battlefield\\n
\"],\\n        \"semantic_type\": \"\",\\n        \"description\": \"\\n
\"},\\n      {\\n        \"column\": \"label\",\\n        \"properties\": {\\n
\"dtype\": \"category\",\\n        \"num_unique_values\": 1,\\n
\"samples\": [\\n          \"Neutral\\n        ],\\n        \"semantic_type\": \"\",
\\n        \"description\": \"\\n      }\\n      {\\n        \"column\": \"reviews\",\\n        \"properties\": {\\n
\"dtype\": \"string\",\\n        \"num_unique_values\": 3,\\n
\"samples\": [\\n          \"The latest The GST Daily! paper.li /
GKConsultants2... Thanks to @ LiquorMarts\\n        ],\\n        \"semantic_type\": \"\",
\\n        \"description\": \"\\n      }\\n    ]\\n  },\\n  \"type\":\"dataframe\"}

dataset.isna().sum()

index      0
company     0
label       0
reviews    686
dtype: int64

dataset=dataset.dropna()
dataset.duplicated().sum()

```

```
np.int64(2340)

dataset=dataset.drop_duplicates()
dataset.company.value_counts()

company
TomClancysRainbowSix           2328
Verizon                          2319
MaddenNFL                        2315
CallOfDuty                       2314
Microsoft                        2304
WorldOfCraft                     2300
NBA2K                            2299
LeagueOfLegends                  2296
TomClancysGhostRecon            2291
Facebook                         2289
ApexLegends                      2278
johnson&johnson                 2257
Battlefield                       2255
Amazon                           2249
CallOfDutyBlackopsColdWar        2242
FIFA                             2238
Dota2                            2225
Overwatch                         2220
Hearthstone                      2219
HomeDepot                        2216
GrandTheftAuto(GTA)              2208
Borderlands                       2205
Xbox(Xseries)                    2201
Google                            2199
Nvidia                            2198
CS-GO                            2195
PlayStation5(PS5)                2183
Fortnite                          2176
Cyberpunk2077                     2175
AssassinsCreed                   2156
RedDeadRedemption(RDR)           2155
PlayerUnknownsBattlegrounds(PUBG) 2150
Name: count, dtype: int64

dataset.info()

<class 'pandas.core.frame.DataFrame'>
Index: 71655 entries, 0 to 74680
Data columns (total 4 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   index     71655 non-null   int64  
 1   company   71655 non-null   object  
 2   label     71655 non-null   object
```

```
3 reviews 71655 non-null object
dtypes: int64(1), object(3)
memory usage: 2.7+ MB

dataset.label.value_counts()

label
Negative      21698
Positive      19712
Neutral       17708
Irrelevant    12537
Name: count, dtype: int64
```

Data Prerocessing

```
def extract_emoji(text_string): # For Extract
    All Emoji From Text
    emoji = []
    for char in text_string:
        if len(char.encode()) >=3 :
            emoji.append(char)
    return " ".join(emoji)

def extract_profile_name(text): # For Extract
    Profile Names From Text
    result=[]
    rev=text.split('@')
    try :
        for i in range(len(rev)-1):
            t=rev[i+1]
            split_name=t.split()
            result.append(f"@{split_name[0]}")
    return result
    except:
        pass

dataset['emoji']=dataset['reviews'].apply(extract_emoji)
# create a table for emojis
dataset['reviews']=dataset['reviews'].apply(lambda x:x.lower())
# convert all reviews to lower case
dataset['natural_review']=dataset['reviews']
# create a column for natural reviews to compare later
dataset['reviews']=dataset['reviews'].apply(lambda
x:re.sub(extract_emoji(x), ' ',x)) # remove emojis from reviews
dataset['profile_name']=dataset['reviews'].apply(extract_profile_name)
# create a table for profile names
dataset['reviews']=dataset['reviews'].apply(lambda x:re.sub('[^a-zA-Z]', ' ',x)) # remove any special characters and
```



```

    "semantic_type": "\",\n                "description": \"\"\n            }\\n        ]\\n    },\"type\":\"dataframe\""
data_for_model=dataset[['reviews','label']]

```

Build Model

```

import torch.nn as nn
import torch
import torch.optim as optim
from torch.utils.data import DataLoader, Dataset
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

import nltk
nltk.download('stopwords')
nltk.download('punkt')
from nltk.corpus import stopwords
stopwords=stopwords.words('english')
data_for_model['reviews']=data_for_model['reviews'].apply(lambda x: ' '.join([word for word in x.split() if word not in stopwords]))

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]  Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]  Unzipping tokenizers/punkt.zip.

PERCENTILE = 97
print( f"{PERCENTILE}th percentile length Reviews:\n{np.percentile([len(x) for x in data_for_model['reviews'].values], PERCENTILE)}" )

97th percentile length Reviews: 193.0

max_len=200
reviews=data_for_model['reviews']
label=data_for_model['label']
label_encoder=LabelEncoder()
label=label_encoder.fit_transform(label)

tokenizer=Tokenizer(num_words=10000)
tokenizer.fit_on_texts(reviews)
tokenized_reviews=tokenizer.texts_to_sequences(reviews)
seq_reviews=pad_sequences(tokenized_reviews,maxlen=max_len)
x_train,x_test,y_train,y_test=train_test_split(seq_reviews,label,test_size=0.25,random_state=42,shuffle=True)

print(reviews.values[0])
print(seq_reviews[0])

```

```
class DataClass(Dataset):
    def __init__(self, review, label):
        self.review=torch.tensor(review)
        self.label=torch.tensor(label)

    def __len__(self):
        return len(self.review)

    def __getitem__(self, idx):
        return self.review[idx], self.label[idx]

def batch(batch):
    texts, labels = zip(*batch)
x_train,x_test,y_train,y_test=train_test_split(seq_reviews,label,test_size=0.25,random_state=42,shuffle=True)

x_train=torch.tensor(x_train)
y_train=torch.tensor(y_train)
x_test=torch.tensor(x_test)
y_test=torch.tensor(y_test)

train_loader=DataLoader(DataClass(x_train,y_train),batch_size=32)
valied_loader=DataLoader(DataClass(x_test,y_test),batch_size=32)

class LSTMClassifier(nn.Module):
    def __init__(self, vocab_size, embed_dim, hidden_dim, output_dim):
        super().__init__()
        self.embedding = nn.Embedding(vocab_size, embed_dim)
        self.lstm = nn.LSTM(embed_dim, hidden_dim, batch_first=True)
        self.fc = nn.Linear(hidden_dim, output_dim)
        self.dropout = nn.Dropout(0.3)

    def forward(self, x):
```

```

        embedded = self.embedding(x)
        _, (hidden, _) = self.lstm(embedded)
        out = self.dropout(hidden[-1])
        return self.fc(out)

vocab=len(tokenizer.index_word)
number_of_class=len(label_encoder.classes_)
model = LSTMClassifier(vocab_size=vocab, embed_dim=100, hidden_dim=64,
output_dim=number_of_class)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Train Model

for epoch in range(10):
    model.train()
    total_loss = 0
    for texts, labels in train_loader:
        optimizer.zero_grad()
        outputs = model(texts)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        total_loss += loss.item()
    print(f'Epoch {epoch+1}, Loss: {total_loss:.4f}')

Epoch 1, Loss: 1855.7926
Epoch 2, Loss: 1276.4628
Epoch 3, Loss: 907.4885
Epoch 4, Loss: 670.0378
Epoch 5, Loss: 527.2335
Epoch 6, Loss: 438.6749
Epoch 7, Loss: 374.1906
Epoch 8, Loss: 333.5186
Epoch 9, Loss: 306.5507
Epoch 10, Loss: 275.7264

# Evaluate the Model
model.eval()
correct, total = 0, 0
with torch.no_grad():
    for texts, labels in valied_loader:
        outputs = model(texts)
        preds = torch.argmax(outputs, dim=1)
        correct += (preds == labels).sum().item()
        total += labels.size(0)
print(f'Validation Accuracy: {correct/total:.2f}')

Validation Accuracy: 0.85

```